

## Hydra development notes

# 1 Minimal compositor

We start with the bare minimum to get a compositor up and running with Mir. This was based on the example code here:

---

```
#include <miral/runner.h>
#include <miral/x11_support.h>
#include <miral/minimal_window_manager.h>

using namespace miral;

int main(int argc, char const* argv[]) {
    MirRunner runner{argc, argv};

    return runner.run_with({
        X11Support{},
        set_window_management_policy<MinimalWindowManager>()
    });
}
```

---

After referencing the more in-depth example, I added `X11Support{}` to the options.

## 1.1 Building

A basic cmake build script based on the example script:

---

```
include(FindPkgConfig)
pkg_check_modules(miral REQUIRED miral IMPORTED_TARGET)

add_executable(hydra main.cpp)
target_link_libraries(hydra PRIVATE PkgConfig::miral)
```

---

### 1.1.1 .deb builds

For consistency, .deb builds are done in a Docker container. The actual .deb build step is generated using CPack:

---

```
install(TARGETS hydra RUNTIME DESTINATION bin)

set(CPACK_PACKAGE_CONTACT jordan@carbondoes.dev)
set(CPACK_DEBIAN_PACKAGE_DEPENDS
    "libmiral5t64"
```

```
"mir-graphics-drivers-desktop")
list(JOIN CPACK_DEBIAN_PACKAGE_DEPENDS ", " CPACK_DEBIAN_PACKAGE_DEPENDS)

include(CPack)
```

---

## 1.2 Launching applications

At this point, the compositor launches, but there is no way (convenient) way to launch an application. Apps can be launched manually, though:

---

```
export WAYLAND_HOST=99
hydra
sleep 30
xfce4-terminal
```

---

## 2 Shell

I decided to work on an application launcher idea I've had for a while: a recreation of Emacs' Hydra. This provides 2 main features:

- Instead of using key combinations where you hit all the keys at once, a string of keys is entered, starting with a leader key that brings up the hydra.
- Commands are organized into a trie, where each node of the trie is visually presented as a table of options.
- Other UI elements (e.g. fuzzy search) can also be displayed to allow for selection from lists

### 2.1 Creating a Window

I chose to use Dear ImGui to build the widgets. ImGui needs a window backend to function, which handles creation of the window surface, keyboard events, etc.

#### 2.1.1 Reusing an existing wl\_display connection

The windowing backend needs to support using an existing display connection in order for the shell to act as an internal client.

( Initially I searched for a way to do this with GLFW, but there are only functions to get the `wl_display` GLFW creates, rather than passing in our own. )

SDL3 has improved wayland support and does support creating an `SDL_Window` with an existing `wl_display`:

---

```
SDLContext::SDLContext(wl_display* display) {
    if(display) {
```

```
    SDL_SetPointerProperty(SDL_GetGlobalProperties(),
        ↪ SDL_PROP_GLOBAL_VIDEO_WAYLAND_WL_DISPLAY_POINTER, display);
}
```

---

Because SDL3 is not yet officially packaged, I chose to bundle it with the project and statically link it:

---

```
set(SDL_STATIC ON CACHE BOOL "Build SDL static target" FORCE)
set(SDL_SHARED OFF CACHE BOOL "Don't build SDL shared target" FORCE)
FetchContent_DeclareEx(
    SDL3
    GIT_REPOSITORY https://github.com/libSDL-org/SDL.git
    GIT_TAG         eb3fc0684c9957de108df5b338307418b0130c3b
    OVERRIDE_FIND_PACKAGE
)
```

---

While SDL3 supports using an existing `wl_display`, but this had a few quirks:

- This bug I reported where the `wl_display` was closed during initialization, even though if it was not created/owned by SDL. Initially I patched this locally, but they fixed it very quickly upstream.
- If the `WAYLAND_DISPLAY` environment variable is not set, init fails, even though we already have a `wl_display`. Setting it to empty string if it's not set is enough:

---

```
// We know we have a display, so avoid failing Init() if the env is not set
SDL_SetEnvironmentVariable(SDL_GetEnvironment(), "WAYLAND_DISPLAY", "",
    ↪ false);
```

---

### 2.1.2 RAII wrappers

SDL is a C library, and expects you to call `SDL_DestroyWindow()`. For this, and a couple other things, I wrote simple RAII wrappers that call the corresponding destroy function as the object exits scope. For simplicity, most of these wrappers are unmovable / uncopyable.

### 2.1.3 Custom window role with `wlr-layer-shell`

At this point, making a normal window with `SDL_CreateWindow()` is simple. However, this creates a normal floating window, but we want to make an overlay.

Referencing Waybar, I found that it uses the `wlr-layer-shell` extension to configure the overlay.

Out of the box SDL creates an `xdg-toplevel-surface` for its windows, unless you set `SDL_PROP_WINDOW_CREATE_WAYLAND_SURFACE_ROLE_CUSTOM_BOOLEAN`, in which case it will create a `wl_surface` without a role.

From the SDL3 wayland README, getting the `wl_surface` from the `SDL_Window` is simple:

---

```
auto* surface = static_cast<wl_surface*>(SDL_GetPointerProperty(window_props,
    SDL_PROP_WINDOW_WAYLAND_SURFACE_POINTER, nullptr));
```

---

Setting up the `wlr-layer-surface` role is a little more involved:

- I attempted to use `wayland-scanner` to generate interface code, but this had some conflicts with C++ keywords, so I ended up using `hyperwayland-scanner`. Referencing KDE's `ecm_add_wayland_client_protocol`, I wrote my a small `cmake` wrapper:

---

```
function(add_wayland_client_protocol target)
    set(options "")
    set(multiValueArgs PROTOCOLS)
    set(oneValueArgs "")
    cmake_parse_arguments(PARSE_ARGV 1 arg "${options}" "${oneValueArgs}"
        ↪ "${multiValueArgs}")

    foreach(protocol ${arg_PROTOCOLS})
        get_filename_component(_protocol ${protocol} ABSOLUTE)
        cmake_path(GET _protocol STEM _protocol_name)

        set(_outputs
            "${CMAKE_CURRENT_BINARY_DIR}/${_protocol_name}.hpp"
            "${CMAKE_CURRENT_BINARY_DIR}/${_protocol_name}.cpp")
        add_custom_command(
            OUTPUT ${_outputs}
            COMMAND $<TARGET_FILE:hyperwayland-scanner> --client ${_protocol}
                ↪ ${CMAKE_CURRENT_BINARY_DIR}
            DEPENDS ${_protocol} VERBATIM
        )
        target_sources(${target} PUBLIC ${_outputs})
    endforeach()
endfunction()
```

---

- The `wl-layer-shell` proxy is obtained from the registry. I was initially skeptical that the registry listener callback was guaranteed to be called after one `wl_display_roundtrip`, but `SDL` appears to do the same thing. This makes it easier to write a wrapper that fetches a single global from the registry:

---

```
RegistryListener(struct wl_display* display, wl_interface const* target, int
    ↪ version)
    : target(target), version(version)
{
    auto* registry = wl_display_get_registry(display);
    wl_registry_add_listener(registry, &global_listener, this);
    wl_display_roundtrip(display);
    wl_registry_destroy(registry);
```

```
}
```

```
// other boiler plate
```

---

- After we get the `LayerShell` we use `get_layer_surface(wl_surface*, ...)` to get a `std::weak_ptr<LayerSurface>` for our `wl_surface`. I chose to have the `LayerShell` own the `LayerSurface`'s, so all of the surfaces are destroyed before the shell object.
- There is some additional configuration for the surface (anchors, `exclusive_zone`, etc.)

I wrapped all of this up in a `LayerWindow` class extending the basic `Window` class.

### 2.1.4 Reacting to fullscreen clients

If another client enters fullscreen, we don't want to draw the shell on top of it (unless we are currently interacting with it).

( Initially, I considered doing this in a `WindowManager`, in `handle_modify_window()`. Whenever a client enters fullscreen, it should generate an event. However, when I tested this I wasn't getting any events. It turns out this was a bug in Mir, only affecting Xwayland clients, which I tracked down later. )

Referencing Waybar again, I found that they used `wlr-foreign-toplevel`, which provides information directly to clients about modifications to other surfaces, including fullscreen state.

I wrote a simple wrapper around it, counting the number of fullscreen toplevels as they are created/destroyed/modified. This is then used in `LayerWindow::should_hide() -> bool`, return true if the count is positive.

## 2.2 Widgets

The Shell UI is broken down into a couple different widgets, described by a couple concepts in `widget.h`

- `IsDrawable`: any object defining `should_draw() -> bool` and `draw()`
  - `StatusLine`: Always visible bar designed to show the window title, date, and error/status messages.
  - `IsPrompt`: drawable with `handle_key(Key)` and `try_result() -> optional<value>`
    - \* `TablePrompt`: A single page of commands
    - \* `SearchPrompt`: A list of options with fuzzy search input

These widgets are designed to be drawn within a `FrameContext`, an object managing the lifetime of each frame, essentially:

1. `ImGui::NewFrame()`
2. Widgets drawn here (generates imgui draw data)
3. `ImGui::EndFrame()`

4. Render the ImGui draw data
5. Swap buffers

To facilitate this, `FrameContext::start_frame()`, does (1), and returns a guard object, which does (3), (4), and (5) on destruction. Optionally, a frame can be marked hidden, skipping (4).

The `TablePrompt` widget is relatively simple, but the `SearchPrompt` needs to filter the list of options based on the inputted text, so I used `rapidfuzz`.

## 2.3 Run loop

---

```
void Shell::run(Window& window, Callback& cb) {
    self->is_done = false;

    FrameContext fc(&window);
    while(!done()) {
        if(auto res = self->frame(window, fc)) {
            cb(res.value());
        }
    }
}
```

---

The main run loop of the Shell takes a `Window`, and a `Callback`:

- The `Window` is not owned by the `Shell` to maintain separation from the window configuration, and I wanted to avoid having to make `Window` moveable/copyable
- A `Callback` is a type-erased lambda where prompt results are returned. Displaying new prompts and other modifications to the `Shell` should be done through this callback for thread-safety reasons, except a few functions that have their own locks (e.g. `handle_key`).

I chose not to use `std::function` primarily for ownership reasons: I did not want to limit the callbacks to copyable lambdas.

1. Each frame starts by handling window events, breaking early if there is a result:

---

```
// Self::frame()
fc.handle_events([this](const SDL_Event& e){
    handle_event(e);
});

if(auto res = pop_result()) {
    return res;
}
```

---

2. Then the widgets are drawn:

---

```
// Self::frame():
auto frame_guard = fc.start_frame();

// Self::draw():
// ... imgui positioning code ...
status.draw();

if(cur_prompt.has_value() && should_draw(*cur_prompt)) {
    // ... imgui positioning code ...
    ::hydra::shell::draw(cur_prompt.value());
}
```

---

## 3 Integrating the shell

### 3.1 Starting an internal client

The Shell state management (starting the internal client, providing a callback, managing the lifetime of the Shell) is wrapped up in the `ShellLauncher` class.

From the mir documentation, internal clients provide 2 callbacks, one accepting the session pointer (the application), and one accepting the `wl_display` pointer.

The callback accepting the `wl_display` will run on its own thread, so this where the Shell's main loop will run:

---

```
void ShellLauncher::operator()(struct wl_display* display) {
    pthread_setname_np(pthread_self(), "HydraShell");

    // wait for the session to be set
    {
        std::unique_lock lock{session_lock};
        startup_cv.wait(lock, [this]{ return !weak_session.expired(); });
    }

    using namespace hydra::shell;

    SDLContext ctx(display);
    LayerWindow window(ctx, Window::Properties::FromConfig());

    auto cb = Shell::Callback::Create([](auto){ /* do nothing */ });
    shell.run(window, cb);
}
```

---

We also need to enable the required wayland extensions first:

---

```

void ShellLauncher::enable_extensions(miral::WaylandExtensions& extensions) {
    for(auto ext: std::array {
        miral::WaylandExtensions::zwlr_layer_shell_v1,
        miral::WaylandExtensions::zwlr_foreign_toplevel_manager_v1,
    }) {
        extensions.conditionally_enable(ext,
            ↪ [this](miral::WaylandExtensions::EnableInfo const& info) {
                if(auto session = weak_session.lock()) {
                    return session == info.app();
                }

                return false;
            });
    }
}

```

---

ShellLauncher isn't copyable, so we wrap it in a lambda before passing it to the runner:

---

```

auto internal_client() {
    return [this](auto&& args) {
        operator()(std::forward<decltype(args)>(args));
    };
}

```

---

## 3.2 StateMachine

At this point the shell is drawn to the screen as we expect, but doesn't accept input. In order to simplify writing a Callback, I wrote a derived class: **StateMachine**.

---

```

template <typename Input>
template <typename T, State... states>
auto StateMachine<Input>::Create(auto&&... args) ->
    ↪ std::shared_ptr<StateMachine>;

```

---

- T... is a base class, defining the states
- states... represents a static map from an index (std::size\_t) a state member function (auto(T::\*)())
- Upon jumping to a new state, the corresponding member function is called and the result becomes the new current state
- When the callback is called, the input is passed to the current state, which returns a new state to jump to, or -1 for no-op.

This allows us to define some basic interaction:



---

```

enum State {
    IDLE,
    COMMAND,
};

auto ShellLauncher::idle() {
    return [](auto) -> std::size_t {
        throw std::runtime_error("Unreachable");
    };
}

auto ShellLauncher::command() {
    enum Commands: Option::value_t {
        QUIT,
    };

    shell.show(hydra::Table{
        std::pair{hydra::Key::Keycode(SDLK_Q),
        ↪ hydra::Option{std::size_t(Commands::QUIT), "Quit"}},
    });

    return [this](auto res) -> std::size_t {
        // ... error handling ...

        switch(Commands(res)) {
            case Commands::QUIT:
                runner->stop();
                break;
            default:
                break;
        }

        return State::IDLE;
    };
}

ShellLauncher::ShellLauncher(MirRunner* runner)
: runner(runner) {
: runner(runner),
    state_machine(StateMachine::Create<
        ShellLauncher,
        hydra::util::State { State::IDLE, &ShellLauncher::idle },
        hydra::util::State { State::COMMAND, &ShellLauncher::command }
        >(this)) { /* ... */ }

bool ShellLauncher::show_commands() {
    bool ret = false;

```

```

state_machine->lock([&ret](std::size_t state) -> std::size_t {
    if(state == State::IDLE) {
        ret = true;
        return State::COMMAND;
    }

    return -1;
});

return ret;
}

```

---

Then we can call `show_commands()` if we receive the leader key to bring up the `command()` state and display the table.

### 3.3 Launching applications

Launching applications is done using the `miral::ExternalClientLauncher`, but it doesn't appear to provide a mechanism for reading desktop entries to get a list of applications.

Initially, I looked to rofi for an example, but they appeared to be parsing the desktop entries manually. I also looked at using gnome-menus but it doesn't really appear to be designed to be used as a standalone library. I briefly considered writing my own parser.

Eventually I stumbled into Gio (part of Gtk's Glib) which provides a very clean interface for this. With Giomm (the C++ wrapper for Gio), getting a list of applications is a one liner:

---

```
std::vector<Glib::RefPtr<Gio::AppInfo>> apps = Gio::AppInfo::get_all();
```

---

After this, launching an application from the list is simple:

---

```
auto cmds = ExternalClientLauncher::split_command(selected->get_commandline());
this->launcher->launch(cmds);
```

---

### 3.4 Window Management

`WindowManager` primarily handles listing windows, getting the focused window, and switching focus. I primarily referenced egmde's `WindowManagerPolicy` for the basic setup. Every new window gets `WindowManager::Metadata` as userdata. The `Metadata` provides a public interface for a window, without having to worry about locks:

- Clients can iterate over the list of windows with `WindowManager::locked_list_windows` which provides a `std::weak_ptr<Metadata>`.
- Other locked functions accept `std::shared_ptr<Metadata>` as input, using the stored id to lookup the `WindowInfo`

There are only 2 workspaces at the moment. One for the shell, and another default workspace. Each workspace has an associated `WorkspaceInfo`.

### 3.4.1 WorkspaceInfo

WorkspaceInfo provides very similar functionality to `miral::ApplicationSelector` (internal). It tracks the list of windows in the workspace, in most-recently focused first order.

In order to make WorkspaceInfo easier to test alone, WorkspaceInfoImpl is templated with the Workspace and Window type, making mocking easy.

When the shell has exclusive focus, it is not possible to switch to a window in the default workspace. To be able to focus windows from shell commands,

### 3.4.2 Inputs, synchronization, and thread-safety

There are 2 different threads interacting with the Shell:

- The main shell thread, in `Shell::run()`
- Compositor keyboard events moving the Shell from `State::IDLE` to `State::COMMAND`

We guard mutation of the state with a lock managed by the `StateMachine`. It also provides a function `StateMachine::lock(func)` to run a function under lock and potentially change the state.

Keyboard events are passed to the `Shell` whenever it is in focus. Pressing the leader key to bring up the command tree, bringing the `Shell` into focus, but it will then also receive the leader key event. To avoid this, keys are instead passed to the `Shell` with `Shell::handle_key(key, timestamp)` which adds the keys to a buffer:

---

```
void handle_key(Key key, uint64_t timestamp_ns) {
    std::lock_guard g{buffer_lock};

    const auto timeout = Config::Get().buffer_timeout<std::chrono::nanoseconds>();
    while(buffer.size() && timestamp_ns - buffer.front().second > timeout) {
        buffer.pop();
    }

    buffer.push({key, timestamp_ns});
    buffer_cv.notify_all();
}
```

---

## 3.5 Command Tree

We don't just want to display one page of commands, but a trie of commands, where each node is either a page or a command to execute, and the edges are keys.

`template <typename Key, typename Value> struct Trie` holds this information.

- `template <typename trie_ptr_t> Trie::base_node_t` is a handle to a single node:
  - The `trie_ptr_t` is either `Trie*` or `Trie const*`. Const-correctness is simplified with deducing this:

---

```
auto& self(this base_node_t self) { return self.parent->nodes[self.idx];
↳ }
```

---

`self()` returns in `Trie const&` when called from `const` members or in `base_node_t<Trie const*>`, or `Trie&` otherwise. This means the non-`const` members of `base_node_t` cannot be called from `base_node_t<Trie const*>`.

- `root()` and `croot()` return mutable and `const` handles to the root node.
- Node values are default constructed, and can be accessed with the dereference operators
- Can be constructed from a `std::tuple` of values

With one `Trie`, we can list all our commands:

---

```
static const hydra::util::Trie<Key, Option> tree {
    std::tuple{Key::Keycode(SDLK_Q), opt("Quit/logout session"),
        std::tuple{Key::Keycode(SDLK_Q), opt("Quit", QUIT)}}},
    std::tuple{Key::Keycode(SDLK_SPACE), opt("Launch application", LAUNCH)},
    std::tuple{Key::Keycode(SDLK_W), opt("Windows"),
        std::tuple{Key::Keycode(SDLK_W), opt("Find window", WINDOW_FIND)},
        std::tuple{Key::Keycode(SDLK_N), opt("Next window", WINDOW_NEXT)},
        std::tuple{Key::Keycode(SDLK_P), opt("Prev window", WINDOW_PREV)},
        std::tuple{Key::Keycode(SDLK_D), opt("Close window", WINDOW_CLOSE)}}
};
```

---

We can use a mutable lambda to keep track of the current node

---

```
return [this, cur=Commands::tree.croot(), status](auto res) mutable -> std::size_t
↳ {
    // ... error handling ...
    if(auto next = cur.try_get(key)) {
        cur = *next;
        switch(cur->value) {
            case Commands::QUIT:
                this->runner->stop();
                return States::IDLE;
            // ... handle other commands ...
            case Commands::NONE:
                show_node(&shell, *cur);
                return -1;
        }
    }
}
```

---